

Leaf-alias: real-time importance sampling of dynamic emissive textures

Jose Aguerre
Universidad de la República
Montevideo, Uruguay
jpaguerre@fing.edu.uy




Figure 1: Two scenes lit by a single dynamic, video-driven emissive texture, rendered with our leaf-alias importance sampler. (Left) The Rover scene — a virtual-production LED wall displaying a looping road/mountain video panorama that lights the vehicle directly, sampled here with BudgetLeafAlias. (Right) The Cinema scene — a theater lit only by the animated sunset footage playing on its own screen, sampled with CodecLeafAlias. Each panel compares, left to right: Random sampling, 2D CDF inversion, our technique, and a converged reference, all at their calibrated real-time budgets (≈ 33 ms).

Abstract

We present a method for fast importance sampling HDR textured lights that allows real-time video-driven lighting. Our method partitions a light’s texture into a budgeted, variance driven quadtree and builds an alias table over the resulting leaves, avoiding the cost of a dense per-frame CDF while still capturing the texture’s overall energy distribution. We present two variants: one that reuses the hierarchical structures already present in the metadata of modern codec video streams, and another that builds a budgeted GPU-friendly quadtree to fit the entire sampling structure in L1 cache. Compared to the dense 2D CDF-inversion baseline, at a matched real-time frame-time budget our implemented technique reduces RMSE by 9–53% depending on scene, letting the saved sampling cost be reinvested in additional light samples. It also reduces memory

consumption and transfers by several orders of magnitude when compared to state of the art samplers.

CCS Concepts

• Computing methodologies → Reflectance modeling.

Keywords

These are the keywords

ACM Reference Format:

Jose Aguerre. 2025. Leaf-alias: real-time importance sampling of dynamic emissive textures. *ACM Trans. Graph.* 1, 1 (September 2025), 2 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 Introduction

Importance sampling of emissive textures is essential for efficient image-based lighting, and increasingly critical in production workflows where video is used directly as a dynamic light source in virtual production and VFX. A single emissive texture lets a renderer reproduce the directional detail of a real environment, far beyond what analytic lights can approximate. Traditionally this texture has been static, captured once and reused for an entire sequence. Increasingly, though, it is video: LED-wall virtual production, AR/VR object insertion, and video-based environment lighting

Author’s Contact Information: Jose Aguerre, Universidad de la República, Montevideo, Uruguay, jpaguerre@fing.edu.uy.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM 1557-7368/2025/9-ART
<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

all require illumination that changes every frame, not a single fixed capture.

Traditional texture light sampling techniques are designed for static images. They rely on the offline building of large sampling structures that are reused for rendering each frame as long as the image doesn't change. The most common example of such strategy is how rendering and ray tracing toolboxes deal with HDR environment maps. The state-of-the-art approach is to build a two-dimensional piecewise constant distribution from the texture's luminance: a marginal CDF over rows and a per-row conditional CDF. The sampler is then inverted via binary search at sample time, which is expensive for large textures. A more recent and well-adopted approach is to generate a mipmap-style hierarchical importance map [1]. At sample time, a hierarchical descent through a mip pyramid of the texture is performed, recursively choosing one of four children at each level in proportion to their summed weight. This technique is also expensive and requires several iterations per sample, although it is more GPU friendly because the number of iterations is always the same for each sample, unlike in the binary search. Moreover, the mipmap can be fastly constructed using the existing dedicated hardware on modern GPUs. Both techniques require at least $O(N)$ memory for the sampler, with N being the number of pixels.

Another less used but known technique for sampling emissive textures is the Alias method [2]. This technique draws from an arbitrary discrete distribution in $O(1)$ time after an $O(N)$ one-time build, at the cost of losing the CDF's implicit spatial coherence: two adjacent table buckets can alias to two arbitrary, unrelated locations in the original distribution. This fact is known to disrupt the stratification or discrepancy-reduction of the input point set. It also requires $O(N)$ memory to store the correlation between every pixel. The process for building the Alias table, despite being $O(N)$, is difficult to optimize for real-time construction [3].

We present a new sampling method that is fast to build (under 1 millisecond for 4k HDR images), produces a sampler that is tiny, and is $O(1)$ to sample without disrupting the stratification of the Quasi Monte Carlo input. The core idea is to reuse the variance-driven image partition that is already present in modern video codecs, which already compute hierarchical spatial partitions that adapt to local signal variance [4]. These structures implicitly encode the spatial distribution of luminance and high-frequency content (see Fig. ??). We also propose a variant that builds the partition tree on a budgeted number of leaves, using a Greedy/best-first tree-structured resource allocation [5]. Both techniques show consistent quality gains when rendering video light textures in real-time (33 milliseconds per frame), since they allow executing more texture samples without introducing stratification error.

2 Previous works

Importance sampling of emissive textures is a well-studied problem in the graphics literature in the context of image-based lighting. Most of the works focus on pre-processing the images to allow for exact or near-exact sampling without introducing bias [6]. Several methods choose to simplify the textured light into a set of point or area lights [7], allowing for real-time sampling and faster convergence but at the cost of more error in the solution. Cline et al. [8]




Figure 2: An example of a variance-driven quadtree built by the VP9 video codec. High variance zones are refined into a dense mesh of small cells.

compare several tabular inversion methods, noting the computational order of building each sampler structure and sampling each point. They analyze the 2D CDF inversion via binary search and the per-pixel Alias method, where they flag the disruption of the stratification of the input point set. Per-pixel Alias's raster-flattened table produces a real, measurable periodic banding pattern (Figure ??) that becomes clear in the sample set's power spectrum. We show that our LeafAlias method overcomes those issues because the partition structure helps with stratification.




Figure 3: 2D power spectrum of samples from 2D CDF inversion, plain per-pixel alias, and our LeafAlias, from the same Halton input. Per-pixel alias shows a sharp, isolated off-axis peak (flatness 0.270, peak/median 16,113) – a periodic banding artifact from its flat array index resonating with the input sequence's base-2 digit structure. Both CDF inversion and ours stay close to the input's own clean spectrum.

Despite being $O(1)$ to sample, the per-pixel Alias method is heavy in memory consumption and in the construction of the tables via Vose's technique. Nevertheless, the Alias method is widely used in mesh lights sampling: geometry with thousands of emissive triangles. Moreau et al. [9] maintain a light-sampling BVH incrementally across frames for thousands of moving lights. Bitterli et al.'s ReSTIR [10] resamples across spatial and temporal neighbors to pick among many dynamic lights at a fraction of the cost of resampling from scratch. Conty Estevez and Kulla [11] adaptively split a light tree per shading point rather than committing to one fixed cut. At that scale of thousands of triangles, a plain alias table over

per-triangle flux is common practice and cheap to rebuild every frame.

Some previous works also address video-driven lighting. Havran et al. [?] and Iorns and Rhee [?] both light scenes from captured video environment maps, but neither targets real-time per-frame resampling of the video itself as the importance distribution. Lu et al. [?] rebuild a tabulated per-cube-face CDF in real-time for sampling of a dynamic environment map, noting it as an expensive part of the process. Kronander et al [?] compute real-time video-based lighting using tabulated CDFs and temporal sample reuse to improve convergence.

Attempts to make the per-frame rebuild itself cheaper on a GPU point the same direction: Morrical and Zellmann [?] accelerate 2D CDF inversion with ray-tracing hardware at high construction cost and non-trivial implementation complexity, and Lehmann et al. [?] study GPU-specific weighted sampling and alias-table construction directly.

Regarding the use of tree structures for sampling, Ergun et al. [?] split an environment map into a kd-tree via the greedy partition strategy, to a predefined block count, then fit the sorted leaf probabilities with a closed-form curve for direct inversion. They report gains in memory compression, higher speed and accuracy than CDF inversion at matched performance for static environment maps. One disadvantage of using kd-trees is that evaluating the pdf of the map direction when sampled from BRDF in MIS requires traversing the tree. Lambers [?] proposes a fixed-grid sampler and notes that replacing the grid with a variance-driven quadtree might improve

their results. Rodrigues et al. [?] train a neural normalizing flow that allows both sampling and evaluating the pdf, showing considerable performance gains at the cost of pre-training the structure for the static map. Wu [?] also preprocess the map by dividing the emissive and non-emissive parts into separate structures.

Current open-source renderers mostly rely on dense CDF or mipmapped samplers. Falcor [?] uses a hierarchical mip-pyramid importance map over the environment texture's luminance and, separately, the triangle-scale alias table for mesh lights. Mitsuba 3 [?] builds a Hierarchical2D warp over the environment map – the same mip-pyramid family as Falcor's. pbrt-v4 [?] uses a dense CDF distribution, inverted via binary search. Blender Cycles [?] builds the same kind of dense marginal/conditional CDF over a downsampled background map. Every one of the four either assumes the environment texture is built once and reused.

3 Method

This is the method

3.1 Codec-leaf-alias

3.2 Budget-leaf-alias

4 Results

These are the results

5 Conclusions and future work

These are the conclusions